



Lab 2 Submission - Database Query Performance Profiler


Backlog with Epics and User Stories




Story Point Assignments




Sprint Board (Active Sprint View)



Burndown Chart - Sprint 1



Burndown Chart - Sprint 2



Reflection

Did your estimations reflect the actual effort?

Mostly. I gave the higher-complexity stories (like the 5,000 records/min ingestion requirement and the composite index recommendation logic) 8 points since they involve more moving parts, and smaller UI-type stories like CSV/PDF export got 2-3 points. Since this was a simulated sprint and not real development, I can't say for sure the points match real effort, but the relative sizing between stories felt reasonable.

Was your backlog well-prioritized?

Yes, I put the core pipeline stories first (log ingestion and index recommendation) as High priority since nothing else works without them, and pushed reporting/export features to Medium and Low since they depend on the pipeline already existing.

How did your simulated sprint align with your plan?

I split the 8 stories into two sprints by dependency - Sprint 1 covered ingestion and index recommendation (Epics 1 and 2), Sprint 2 covered the digest and export features (Epics 3 and 4). Since I simulated both sprints quickly instead of over a real week, all stories moved to Done almost immediately rather than being spread out.

What insights did the burndown chart give about your team's capacity?

Because I completed all the work at once instead of gradually, the burndown line dropped straight to zero instead of following the grey guideline. In a real sprint that pattern would suggest the team finished way early or the work was overestimated, so the team could probably take on more stories per sprint. It also shows why the guideline line matters - it's the reference point for knowing if progress is on track.